

Improving QSAR Modeling for Predictive Toxicology using Publicly Aggregated Semantic Graph Data and Graph Neural Networks

Joseph D. Romano*, Yun Hao*, and Jason H. Moore†

*Institute for Biomedical Informatics, University of Pennsylvania,
Philadelphia, Pennsylvania 19104, United States*

†Corresponding author e-mail: jhmoore@upenn.edu

**These authors contributed equally.*

Quantitative Structure-Activity Relationship (QSAR) modeling is the most common computational technique for predicting chemical toxicity, but a lack of methodological innovations in QSAR have led to underwhelming performance. We show that contemporary QSAR modeling for predictive toxicology can be substantially improved by incorporating semantic graph data aggregated from open-access public databases, and analyzing those data in the context of graph neural networks (GNNs). Furthermore, we introspect the GNNs to demonstrate how they can lead to more interpretable applications of QSAR, and use ablation analysis to explore the contribution of different data elements to the final models’ performance.

Keywords: Toxicology; Graph neural networks; Data aggregation; QSAR; Artificial intelligence.

1. Introduction

Evaluating the toxicity of chemicals is an essential component of pharmaceutical and environmental research. Traditionally, the task of establishing toxicity has been accomplished using *in vivo* models, where a model organism is exposed to a chemical of interest and observed for toxic effects, or by performing epidemiological studies on human populations. Both of these approaches are costly and time consuming, and given the hundreds of thousands of compounds of toxicological interest, innovative alternatives are needed to rapidly screen chemicals. In recent decades, predictive toxicology and large-scale chemical screening efforts have emerged to address this issue.

Quantitative Structure-Activity Relationship (QSAR) modeling is arguably the most prevalent method for predicting *in silico* whether a chemical will cause a toxic response. Briefly, QSAR modeling involves collecting regularly structured quantitative descriptions of molecular structures (known as *fingerprints*), and then fitting a statistical model (e.g., logistic regression, or other supervised machine learning algorithms) to sets of chemicals where a toxic endpoint of interest is already known. Since each data point used to train a model is itself the outcome of a single experiment, QSAR is a meta-analysis approach that is complicated not only by the challenge of capturing relevant structural features of chemicals, but also by errors, biases, and ambiguities in the underlying experiments used to generate the training data. Consequently, QSAR is heavily criticized for its disappointing performance on many tasks. The computational toxicology community has long acknowledged the need for new methodological innovations to improve QSAR performance, but few have emerged in the past several decades.

In this study, we address these issues by augmenting the traditional QSAR approach with multimodal graph data aggregated from several public data sources, and analyzing those data in the context of a heterogeneous graph convolutional neural network (GCNN) model. We evaluate the model using 52 assays and their accompanying chemical screening data from the Tox21 data repository, and compare its performance to two rigorously defined traditional QSAR models consisting of random forest and gradient boosting classifiers. Our results show that the GNN strategy significantly outperforms traditional QSAR. We further refine our results by using an ablation analysis to explain the relative contributions of different data sources to the GNNs’ increased performance. Finally, we discuss how GNNs improve the interpretability of QSAR, and suggest future directions to continue this body of work.

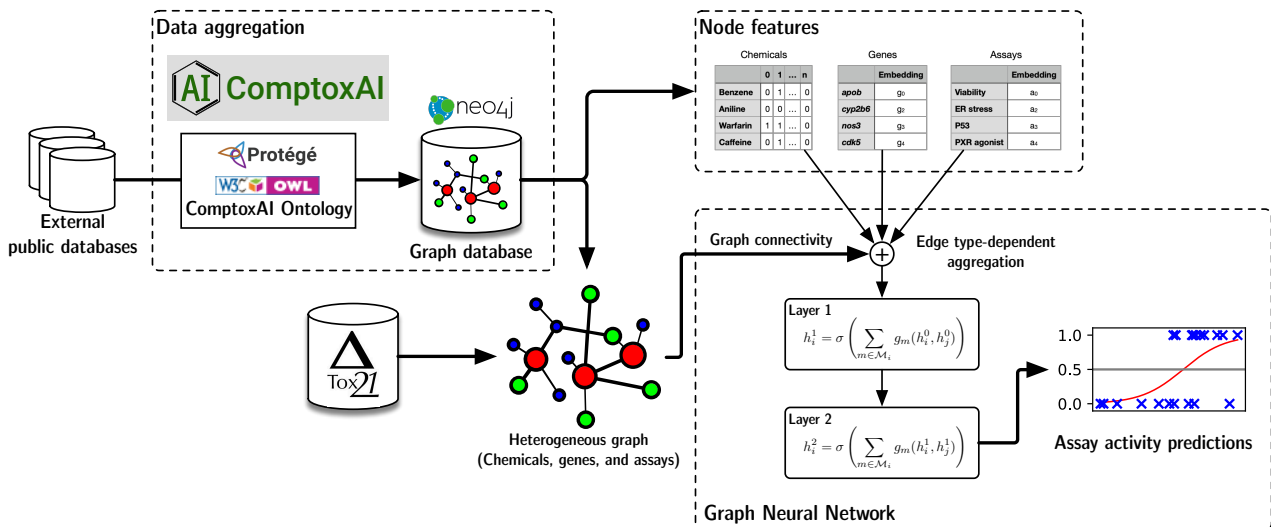


Fig. 1. Schematic overview of the graph machine learning approach used in this study. We build a toxicology-focused graph database (named ComptoxAI) using data aggregated from many public databases, and extract a subgraph for QSAR analysis containing chemicals, assays, and genes. We then train and evaluate a graph neural network that predicts whether or not a chemical activates specific toxicology-focused assays from the Tox21 database.

2. Methods

2.1. Obtaining toxicology assay data

We used the Tox21 dataset—which is a freely available resource produced collaboratively by the US National Institutes of Health, the US Food and Drug Administration, and the US Environmental Protection Agency—to obtain a set of candidate assays for classification and establish ‘ground truth’ relationships between specific chemicals and those assays. Each assay is presented along with experimental screening results describing the response of the assay to exposure to specific chemicals of toxicological interest, which includes pharmaceutical drugs, small molecule metabolites, environmental toxicants and others. We removed all screening experiments with inconclusive or ambiguous results, as well as assays with very few (e.g., < 100) active chemicals.

2.2. Aggregating publicly available multimodal graph data

The graph data used in this study come from a new data resource for computational toxicology, named ComptoxAI. ComptoxAI includes a large graph database containing many entity and relationship types that pertain to translational mechanisms of toxicity, all of which are sourced from third-party public databases (including PubChem, Drugbank, the US EPA’s Computational Toxicology Dashboard, NCBI Gene, and many others). We extracted the subgraph from ComptoxAI’s graph database defined as all nodes representing chemicals, genes, and toxicological assays, as well as the complete set of edges linking nodes of those types. The full graph database for ComptoxAI can be found at <https://comptox.ai>, and will be described in a separate, upcoming publication.

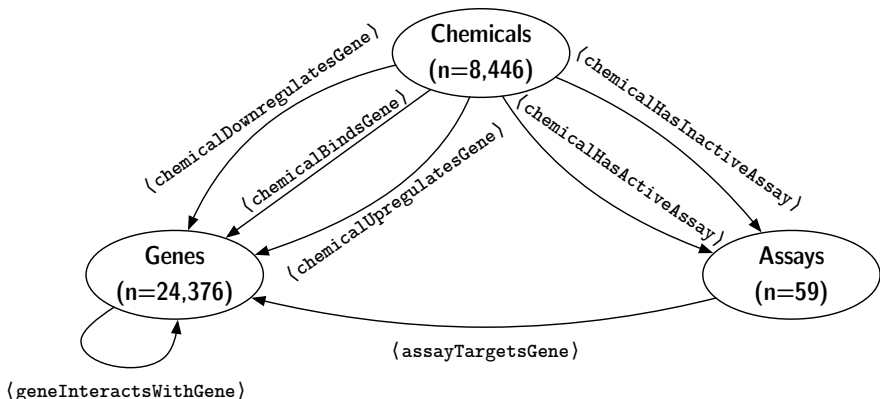


Fig. 2. Metagraph describing the node types, node counts, and edge types in the heterogeneous graph. During implementation of the GNN, we also define corresponding inverse edges (e.g., assay-TargetsGene \leftrightarrow geneTargetedByAssay) to facilitate the message-passing paradigm of the GNN.

The 3 entity types that comprise the nodes of the extracted subgraph are *chemicals*, *assays*, and *genes*. We sourced the chemicals from the US EPA’s DSSTox database, and further filtered them so that each one is equivalent to a distinct compound in PubChem. We obtained genes from the NCBI Gene database, and assays from the Tox21 screening repository as described above. All edges in the graph were sourced from either the Hetionet database or from assay-chemical annotations in Tox21. A metagraph describing the node and edge types in the subgraph is shown in **Figure 2**.

2.3. Heterogeneous graph neural network

We constructed a heterogeneous graph convolutional neural network (GCNN) architecture for the graph ML experiments. Since our graph contains multiple entity types (chemicals, genes, and assays)—each with possibly different sets of node features, and linked by multiple semantically distinct edge types—the architecture extends the common GCNN model to learn separate message passing functions for each edge type. Briefly, each layer of the network aggregates signals from adjacent nodes in the graph, such that a greater number of layers results in signals being aggregated from an increasingly wider radius around each node. The

output of the network can be thought of as encoded representations of nodes that incorporate information from the other nodes in their local neighborhood. The GCNN can be thought of as a generalization of convolutional neural networks (CNNs) used in computer vision—instead of the convolutional operator aggregating signals from nearby pixels in an image, it aggregates features from adjacent nodes in the graph.

In a heterogeneous graph, different node types represent different types of entities, each represented within a semantically distinct feature space. Therefore, the process of aggregating information from adjacent nodes must take those nodes’ types into account. Additionally, different edge types (e.g., ‘chemicalUpregulatesGene’ and ‘chemicalDownregulatesGene’) convey their own semantically distinct meanings, which can substantially effect the flow of information through the network. To handle these two challenges, we learn separate aggregation functions for each edge type in the graph, following the example proposed by Schlichtkrull *et al* in R-GCNs (Relational Graph Convolutional Networks).¹ Within the R-GCN paradigm, the aggregation process can be split into 3 sequential steps: (1.) collecting signals from adjacent nodes using an appropriate edge type-specific message function ϕ , (2.) combining each of those incoming signals (across all edge types) via a reduce function ρ , and (3.) finally updating the target node v by applying an update function ψ . Training the network is roughly equivalent to finding an appropriate parameterization of ϕ for each edge type.

A formal description of the GNN is given in **Appendix A**.

2.3.1. Node classification

Given the GCNN architecture described above, we construct a heterogeneous graph where chemicals are labeled according to whether they do (1) or do not (2) activate an assay of interest. Although we remove the node representing the assay of interest*, all other Tox21 assays are included in the graph, and edges between chemicals and the other assays can therefore be used to improve the inferential capacity of the model beyond those of the baseline QSAR models, which only have access to chemical structure. Similarly, interactions involving genes further increase the information available to the model. Finally, the same MACCS chemical descriptors used to train the baseline QSAR models (see below) are included as node features. The procedure we use for labeling the graph is shown in **Algorithm 1**.

The resulting graph \mathcal{G}_a^* containing labeled chemicals is then applied in a supervised learning context, where the model is trained to predict the chemical labels. We use an 80%/20% train/test split on the labeled chemicals, optimize the GCNN’s parameters using the Adam algorithm (a computationally efficient variant of stochastic gradient descent suitable for sparse gradients),² and compute the error between predicted and true labels via cross entropy loss.

To assess the contribution of the MACCS molecular descriptors when added to the GNN, we trained the node classification model both in the presence and in the absence of MACCS bitstrings applied as node features to each chemical. Intuitively, the model trained without MACCS node features performs inference using only the graph’s topological structure—including gene interactions and activity annotations to the other Tox21 assays—while the one

*To prevent information leakage

Algorithm 1 Labeled heterogeneous graph construction for toxicity assay QSAR model.

Let \mathcal{G} be a heterogeneous graph for QSAR, $a \in \mathcal{A}$ be an assay of interest, and $l(c)$ denote an activity label for chemical $c \in \mathcal{C}$ w.r.t. assay a .

for each chemical $c \in \mathcal{C}$ **do**

if \exists an edge (c, r, a) s.t. the edge type of r is `<chemicalHasActiveAssay>` **then**

$l(c) \leftarrow 1$

else if \exists an edge (c, r, a) s.t. the edge type of r is `<chemicalHasInactiveAssay>` **then**

$l(c) \leftarrow 0$

else

$l(c) \leftarrow$ is undefined

\triangleright No label available for node c

end if

end for

$\mathcal{G}_a^* \leftarrow \mathcal{G} \setminus a$

\triangleright Delete node a from the graph to prevent information leakage

return \mathcal{G}_a^*

with MACCS node features additionally has access to the same chemical structure information used in the non-graph (baseline) QSAR models.

Additional details on the node classification approach are given in **Appendix B**.

2.4. Baseline QSAR classifiers

To assess the relative performance of the GNN classification model, we built 2 additional (non-NN) QSAR models that represent rigorously defined benchmarks consistent with current practice in predictive toxicology: A random forest classifier, and a gradient boosting classifier. Each model was trained on the aforementioned MACCS fingerprints of chemicals computed from SMILES strings, with an 80%/20% training/testing split. We tuned 6 hyperparameters for each random forest model, and 5 for each gradient boosting model, as described in **Table S1**. These were tuned using grid search, where the optimal hyperparameter set is defined as the one that minimized binary cross entropy between predicted labels and true labels on the training data.

3. Results

3.1. GNN node classification performance vs. baseline QSAR models

Of the 68 total assays in the Tox21 database, we retained 52 for analysis in the QSAR experiments. The remaining 16 assays were not used due to either a low number of active chemicals or underrepresentation of screened chemicals in the ComptoxAI graph database. Additionally, we discarded compound labels for chemicals with inconclusive or ambiguous screening results.

As shown in **Figure 3**, the GNN model significantly (1-sided p -value [XXX]; [XXX hypothesis test]) outperforms both the random forest and gradient boosting models in terms of area under the receiver operating characteristic curve (AUROC), attaining a mean AUROC of [XXX] (compared to [XXX] and [XXX], respectively). This is robust evidence that the GNN model tends to substantially outperform ‘traditional’ QSAR models. A notable characteristic

of the GNN AUROCs is that their distribution has a higher variance than either the random forest or gradient boosting AUROCs. Anecdotally, this is likely due to diminished sensitivity of the GNN model when trained on assays with few positive examples—neural networks tend to struggle as data become more sparse, which seems to be the case here. We also compared F1-score distributions between the 3 model types; however, the differences between the 3 models are not statistically significant. The relatively low F1-scores in the 3 model types is a result of the class imbalance in the QSAR toxicity assays—all of the assays contain far more negative samples (assay is inactive) than positive samples (assay is active), which results in any false negatives having a magnified impact on F1. The same increased variance observed in GNN model AUROCs is shown in the GNN F1-scores.

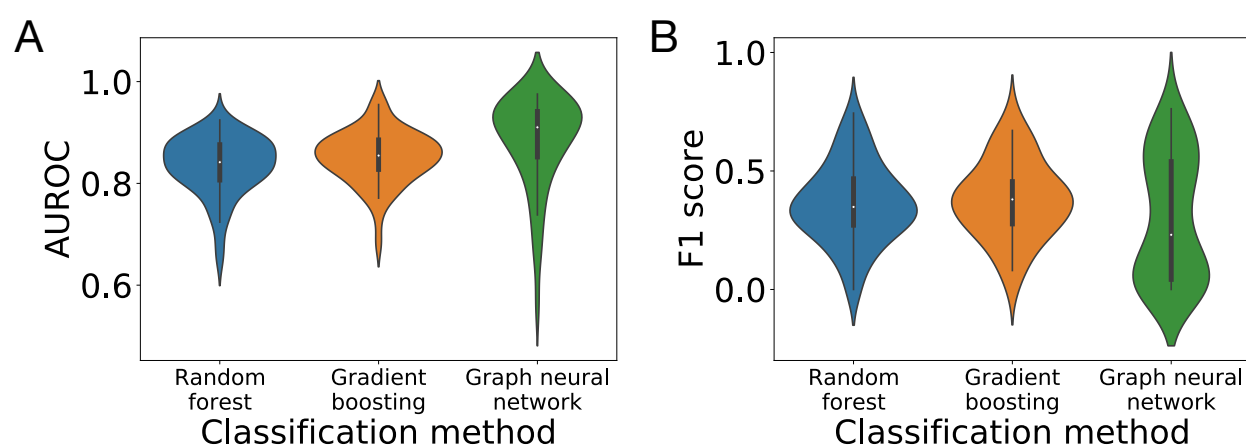


Fig. 3. Overall performance metrics of the 3 QSAR model types on each of the Tox21 assays—a.) area under the receiver operating characteristic curve (AUROC) and b.) F1 score. The mean AUROC is significantly higher for the GNN model than for either of the two baseline QSAR approaches. The differences in F1 scores are not statistically significant. The GNN achieves poor F1 scores on assays with relatively few (e.g., < 100) “active” annotations in Tox21, which is consistent with known performance of neural networks on data with sparse labels.

We performed further review of model performance on two selected assays of interest: PXR agonism (labeled `tox21-pxr-p1` in Tox21) and HepG2 cell viability (`tox21-rt-viability-hepg2-p1`). We selected these assays for two reasons: (1.) Both are semantically distinct from all other Tox21 assays (i.e., there are no other assays measuring pregnane X activity or cell viability), and therefore we would not expect an information leak from other highly correlated Tox21 assays present in the GNN, and (2.) both have a sufficient number of positive chemicals such that their ROC curves attain high resolution at all values of the decision rule across the 3 model types. **Figure 4** shows that the GNN outperforms the random forest and gradient boosting models at virtually all discrimination thresholds in both cases. The high performance of the GNN on HepG2 cell viability is especially noteworthy—cell viability is notoriously challenging to predict in chemical screening experiments. Many of the other 50 Tox21 assays showed similar patterns in performance. All ROC plots are available in

Supplemental Materials.

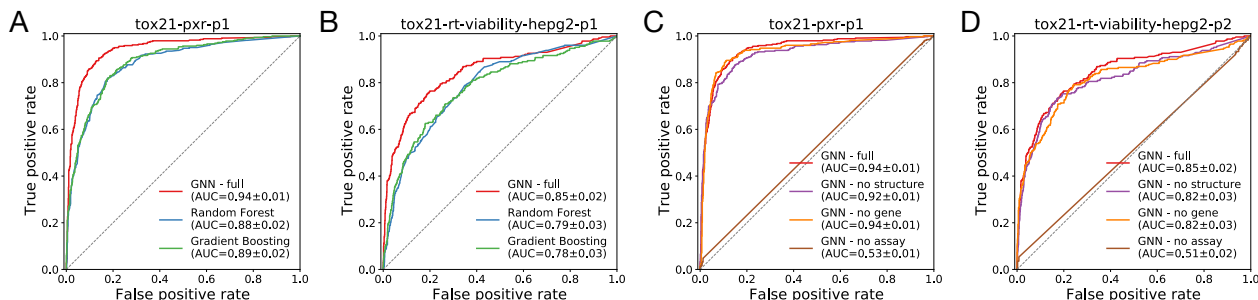


Fig. 4. Receiver operating characteristic (ROC) curves for two selected Tox21 assays: a.) PXR agonism (tox21-prx-p1) and b.) HepG2 cell viability (tox21-rt-viability-hepg2-p1). In both cases, the area under the curve (AUC) is significantly higher for the GNN model than either the Random Forest or Gradient Boosting models.

3.2. Ablation analysis of graph components’ influence on the trained predictive model

To better understand how the GNN model outperforms the random forest and gradient boosting models, we performed an ablation analysis on the two previously mentioned assays—pregnane X agonism and HepG2 cell viability. For both of the assays, we re-trained the model after removing specific components from the GNN:

- All assay nodes.
- All gene nodes.
- MACCS fingerprints for chemical nodes (replacing them with dummy variables so the structure of the network would remain the same).

The results of these experiments are shown in **Figure 4**.

4. Discussion

4.1. GNNs versus traditional ML for QSAR modeling

The toxicology community largely agrees that QSAR underperforms on many tasks, and that new methodological advances are desperately needed. In this study, we demonstrate that GNNs significantly outperform the current gold-standard techniques in the field. Aside from the fact that neural networks can more easily adapt to nonlinear objectives than non-neural network models, this is likely a natural consequence of incorporating biomedical knowledge that goes beyond chemical structure characteristics. Gene interactions provide clues about how chemicals influence metabolic and signaling pathways *in vivo*, and non-target assays (i.e., other assays in the graph aside from the one currently being predicted) may correlate with activity of the target assay.

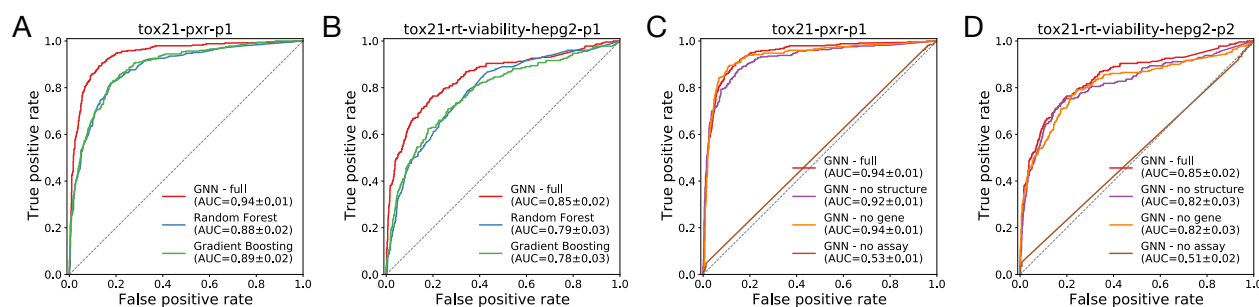


Fig. 5. Receiver Operator Characteristic (ROC) curves for two selected Tox21 assays using different configurations of the GNN model. ‘GNN - full’ is the complete model as described in §2.3.1. ‘GNN - no structure’ omits the MACCS chemical descriptors and replaces them with node embeddings of the same dimensionality. ‘GNN - no gene’ removes gene nodes and their incident edges from the network. ‘GNN - no assay’ removes all assay nodes and incident edges, so predictions are made solely using chemicals, genes, the remaining edges, and the MACCS chemical descriptors as chemical node features.

4.2. Interpretability of GNNs in QSAR

4.3. Future opportunities

5. Conclusions

6. Code availability

All source code pertaining to this study is available on GitHub at <https://github.com/EpistasisLab/qsar-gnn>. A frozen copy of the code at the time of writing is available at [XXX; Zenodo].

7. Supplemental Materials

Supplemental tables and figures are available on FigShare at [XXX].

Acknowledgements

This work was made possible with support from US National Institutes of Health grants R01-LM010098, R01-LM012601, R01-AI116794, UL1-TR001878, UC4-DK112217 (PI: Jason Moore), T32-ES019851, and P30-ES013508 (PI: Trevor Penning).

Appendix A. Graph convolutional network architecture

Some stuff here.

Appendix B. Node classification model

Some more stuff.

References

1. M. Schlichtkrull, T. N. Kipf, P. Bloem, R. Van Den Berg, I. Titov and M. Welling, Modeling relational data with graph convolutional networks, 593 (2018).

2. D. P. Kingma and J. Ba, Adam: A method for stochastic optimization, *arXiv preprint arXiv:1412.6980* (2014).